

Generalizing the convolution operator to irregular domains is typically expressed as a *neighborhood aggregation* or *message passing* scheme. With  $\mathbf{x}_i^{(k-1)} \in \mathbb{R}^F$  denoting node features of node  $i$  in layer  $(k - 1)$  and  $\mathbf{e}_{j,i} \in \mathbb{R}^D$  denoting (optional) edge features from node  $j$  to node  $i$ , message passing graph neural networks can be described as

$$\mathbf{x}_i^{(k)} = \gamma^{(k)} \left( \mathbf{x}_i^{(k-1)}, \bigoplus_{j \in \mathcal{N}(i)} \phi^{(k)} \left( \mathbf{x}_i^{(k-1)}, \mathbf{x}_j^{(k-1)}, \mathbf{e}_{j,i} \right) \right),$$

where  $\bigoplus$  denotes a differentiable, permutation invariant function, e.g., sum, mean or max, and  $\gamma$  and  $\phi$  denote differentiable functions such as MLPs (Multi Layer Perceptrons).

- [The “MessagePassing” Base Class](#)
- [Implementing the GCN Layer](#)
- [Implementing the Edge Convolution](#)
- [Exercises](#)

## The “MessagePassing” Base Class

PyG provides the `MessagePassing` base class, which helps in creating such kinds of message passing graph neural networks by automatically taking care of message propagation. The user only has to define the functions  $\phi$ , i.e. `message()`, and  $\gamma$ , i.e. `update()`, as well as the aggregation scheme to use, i.e. `aggr="add"`, `aggr="mean"` or `aggr="max"`.

This is done with the help of the following methods:

- `MessagePassing(aggr="add", flow="source_to_target", node_dim=-2)` : Defines the aggregation scheme to use ( "add" , "mean" or "max" ) and the flow direction of message passing (either "source\_to\_target" or "target\_to\_source" ). Furthermore, the `node_dim` attribute indicates along which axis to propagate.
- `MessagePassing.propagate(edge_index, size=None, **kwargs)` : The initial call to start propagating messages. Takes in the edge indices and all additional data which is needed to construct messages and to update node embeddings. Note that `propagate()` is not limited to exchanging messages in square adjacency matrices of shape [N N1 only, but can also exchange messages in general sparse assignment matrices, e.g.,  `latest` size=(N, M) as an additional argument. If set to `None`, the

assignment matrix is assumed to be a square matrix. For bipartite graphs with two independent sets of nodes and indices, and each set holding its own information, this split can be marked by passing the information as a tuple, e.g. `x=(x_N, x_M)` .

- `MessagePassing.message(...)` : Constructs messages to node  $i$  in analogy to  $\phi$  for each edge  $(j, i) \in \mathcal{E}$  if `flow="source_to_target"` and  $(i, j) \in \mathcal{E}$  if `flow="target_to_source"` . Can take any argument which was initially passed to `propagate()` . In addition, tensors passed to `propagate()` can be mapped to the respective nodes  $i$  and  $j$  by appending `_i` or `_j` to the variable name, e.g. `x_i` and `x_j` . Note that we generally refer to  $i$  as the central nodes that aggregates information, and refer to  $j$  as the neighboring nodes, since this is the most common notation.
- `MessagePassing.update(aggr_out, ...)` : Updates node embeddings in analogy to  $\gamma$  for each node  $i \in \mathcal{V}$ . Takes in the output of aggregation as first argument and any argument which was initially passed to `propagate()` .

Let us verify this by re-implementing two popular GNN variants, the [GCN layer from Kipf and Welling](#) and the [EdgeConv layer from Wang et al..](#)

## Implementing the GCN Layer

The [GCN layer](#) is mathematically defined as

$$\mathbf{x}_i^{(k)} = \sum_{j \in \mathcal{N}(i) \cup \{i\}} \frac{1}{\sqrt{\deg(i)} \cdot \sqrt{\deg(j)}} \cdot \left( \mathbf{W}^\top \cdot \mathbf{x}_j^{(k-1)} \right) + \mathbf{b},$$

where neighboring node features are first transformed by a weight matrix  $\mathbf{W}$ , normalized by their degree, and finally summed up. Lastly, we apply the bias vector  $\mathbf{b}$  to the aggregated output. This formula can be divided into the following steps:

1. Add self-loops to the adjacency matrix.
2. Linearly transform node feature matrix.
3. Compute normalization coefficients.
4. Normalize node features in  $\phi$ .
5. Sum up neighboring node features ( "add" aggregation).
6. Apply a final bias vector.

Steps 1-3 are typically computed before message passing takes place. Steps 4-5 can be easily processed using the `MessagePassing` base class. The full layer implementation is shown below:

```

import torch
from torch.nn import Linear, Parameter
from torch_geometric.nn import MessagePassing
from torch_geometric.utils import add_self_loops, degree

class GCNConv(MessagePassing):
    def __init__(self, in_channels, out_channels):
        super().__init__(aggr='add') # "Add" aggregation (Step 5).
        self.lin = Linear(in_channels, out_channels, bias=False)
        self.bias = Parameter(torch.empty(out_channels))

        self.reset_parameters()

    def reset_parameters(self):
        self.lin.reset_parameters()
        self.bias.data.zero_()

    def forward(self, x, edge_index):
        # x has shape [N, in_channels]
        # edge_index has shape [2, E]

        # Step 1: Add self-loops to the adjacency matrix.
        edge_index, _ = add_self_loops(edge_index, num_nodes=x.size(0))

        # Step 2: Linearly transform node feature matrix.
        x = self.lin(x)

        # Step 3: Compute normalization.
        row, col = edge_index
        deg = degree(col, x.size(0), dtype=x.dtype)
        deg_inv_sqrt = deg.pow(-0.5)
        deg_inv_sqrt[deg_inv_sqrt == float('inf')] = 0
        norm = deg_inv_sqrt[row] * deg_inv_sqrt[col]

        # Step 4-5: Start propagating messages.
        out = self.propagate(edge_index, x=x, norm=norm)

        # Step 6: Apply a final bias vector.
        out = out + self.bias

    return out

    def message(self, x_j, norm):
        # x_j has shape [E, out_channels]

        # Step 4: Normalize node features.
        return norm.view(-1, 1) * x_j

```

`GCNConv` inherits from `MessagePassing` with "add" propagation. All the logic of the layer takes place in its `forward()` method. Here, we first add self-loops to our edge indices using the `torch_geometric.utils.add_self_loops()` function (step 1), as well as linearly transform node features by calling the `torch.nn.Linear` instance (step 2).

The normalization coefficients are derived by the node degrees  $\text{deg}(i)$  for each node  $i$  which gets transformed to  $1/(\sqrt{\text{deg}(i)} \cdot \sqrt{\text{deg}(j)})$  for each edge  $(j, i) \in \mathcal{E}$ . The result is saved in the tensor `norm` of shape `[num_edges, ]` (step 3).

We then call `propagate()`, which internally calls `message()`, `aggregate()` and `update()`. We pass the node embeddings `x` and the normalization coefficients `norm` as additional arguments for message propagation.

In the `message()` function, we need to normalize the neighboring node features `x_j` by `norm`. Here, `x_j` denotes a *lifted* tensor, which contains the source node features of each edge, *i.e.*, the neighbors of each node. Node features can be automatically lifted by appending `_i` or `_j` to the variable name. In fact, any tensor can be converted this way, as long as they hold source or destination node features.

That is all that it takes to create a simple message passing layer. You can use this layer as a building block for deep architectures. Initializing and calling it is straightforward:

```
conv = GCNConv(16, 32)
x = conv(x, edge_index)
```

## Implementing the Edge Convolution

The [edge convolutional layer](#) processes graphs or point clouds and is mathematically defined as

$$\mathbf{x}_i^{(k)} = \max_{j \in \mathcal{N}(i)} h_{\Theta} \left( \mathbf{x}_i^{(k-1)}, \mathbf{x}_j^{(k-1)} - \mathbf{x}_i^{(k-1)} \right),$$

where  $h_{\Theta}$  denotes an MLP. In analogy to the GCN layer, we can use the `MessagePassing` class to implement this layer, this time using the "max" aggregation:

```

import torch
from torch.nn import Sequential as Seq, Linear, ReLU
from torch_geometric.nn import MessagePassing

class EdgeConv(MessagePassing):
    def __init__(self, in_channels, out_channels):
        super().__init__(aggr='max') # "Max" aggregation.
        self.mlp = Seq(Linear(2 * in_channels, out_channels),
                      ReLU(),
                      Linear(out_channels, out_channels))

    def forward(self, x, edge_index):
        # x has shape [N, in_channels]
        # edge_index has shape [2, E]

        return self.propagate(edge_index, x=x)

    def message(self, x_i, x_j):
        # x_i has shape [E, in_channels]
        # x_j has shape [E, in_channels]

        tmp = torch.cat([x_i, x_j - x_i], dim=1) # tmp has shape [E, 2 * in_channels]
        return self.mlp(tmp)

```

Inside the `message()` function, we use `self.mlp` to transform both the target node features `x_i` and the relative source node features `x_j - x_i` for each edge  $(j, i) \in \mathcal{E}$ .

The edge convolution is actually a dynamic convolution, which recomputes the graph for each layer using nearest neighbors in the feature space. Luckily, PyG comes with a GPU accelerated batch-wise k-NN graph generation method named `torch_geometric.nn.pool.knn_graph()`:

```

from torch_geometric.nn import knn_graph

class DynamicEdgeConv(EdgeConv):
    def __init__(self, in_channels, out_channels, k=6):
        super().__init__(in_channels, out_channels)
        self.k = k

    def forward(self, x, batch=None):
        edge_index = knn_graph(x, self.k, batch, loop=False, flow=self.flow)
        return super().forward(x, edge_index)

```

Here, `knn_graph()` computes a nearest neighbor graph, which is further used to call the `forward()` method of `EdgeConv`.



This leaves us with a clean interface for initializing and calling this layer:

```
conv = DynamicEdgeConv(3, 128, k=6)
x = conv(x, batch)
```

## Exercises

Imagine we are given the following `Data` object:

```
import torch
from torch_geometric.data import Data

edge_index = torch.tensor([[0, 1],
                          [1, 0],
                          [1, 2],
                          [2, 1]], dtype=torch.long)
x = torch.tensor([-1, 0, 1], dtype=torch.float)

data = Data(x=x, edge_index=edge_index.t().contiguous())
```

Try to answer the following questions related to `GCNConv` :

1. What information does `row` and `col` hold?
2. What does `degree()` do?
3. Why do we use `degree(col, ...)` rather than `degree(row, ...)` ?
4. What does `deg_inv_sqrt[col]` and `deg_inv_sqrt[row]` do?
5. What information does `x_j` hold in the `message()` function? If `self.lin` denotes the identity function, what is the exact content of `x_j` ?
6. Add an `update()` function to `GCNConv` that adds transformed central node features to the aggregated output.

Try to answer the following questions related to `EdgeConv` :

1. What is `x_i` and `x_j - x_i` ?
2. What does `torch.cat([x_i, x_j - x_i], dim=1)` do? Why `dim = 1` ?